

u-net-v2

November 8, 2024

U-NET V2 ARCHITECTURE BUILDING

LOADING MODULES

```
[8]: import os
import time
import random
import pathlib
import itertools
from glob import glob
from tqdm import tqdm_notebook, trange

# import data handling tools
import cv2
import numpy as np
import pandas as pd
import seaborn as sns
sns.set_style('darkgrid')
import matplotlib.pyplot as plt
%matplotlib inline
from skimage.color import rgb2gray
from skimage.morphology import label
from skimage.transform import resize
from sklearn.model_selection import train_test_split
from skimage.io import imread, imshow, concatenate_images

# import Deep learning Libraries
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import backend as K
from tensorflow.keras.models import Model, load_model, save_model
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import Adam, Adamax
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from tensorflow.keras.layers import Input, Activation, BatchNormalization, \
↳ Dropout, Lambda, Conv2D, Conv2DTranspose, MaxPooling2D, concatenate

# Ignore Warnings
```

```
import warnings
warnings.filterwarnings("ignore")

print ('modules loaded')
```

modules loaded

SPLITTING THE DATAFRAME

```
[9]: # function to create dataframe
def create_df(data_dir):
    images_paths = []
    masks_paths = glob(f'{data_dir}/*/*_mask*')

    for i in masks_paths:
        images_paths.append(i.replace('_mask', ''))

    df = pd.DataFrame(data= {'images_paths': images_paths, 'masks_paths':
↪masks_paths})

    return df

# Function to split dataframe into train, valid, test
def split_df(df):
    # create train_df
    train_df, dummy_df = train_test_split(df, train_size= 0.8)

    # create valid_df and test_df
    valid_df, test_df = train_test_split(dummy_df, train_size= 0.5)

    return train_df, valid_df, test_df
```

```
[10]: def create_gens(df, aug_dict):
    img_size = (256, 256)
    batch_size = 40

    img_gen = ImageDataGenerator(**aug_dict)
    msk_gen = ImageDataGenerator(**aug_dict)

    # Create general generator
    image_gen = img_gen.flow_from_dataframe(df, x_col='images_paths',
↪class_mode=None, color_mode='rgb', target_size=img_size,
                                         batch_size=batch_size,
↪save_to_dir=None, save_prefix='image', seed=1)

    mask_gen = msk_gen.flow_from_dataframe(df, x_col='masks_paths',
↪class_mode=None, color_mode='grayscale', target_size=img_size,
```

```

                                batch_size=batch_size,
↪save_to_dir=None, save_prefix= 'mask', seed=1)

gen = zip(image_gen, mask_gen)

for (img, msk) in gen:
    img = img / 255
    msk = msk / 255
    msk[msk > 0.5] = 1
    msk[msk <= 0.5] = 0

    yield (img, msk)

```

ARCHITECTURE BUILDING

```

[11]: def ChannelAttention(x, ratio=8):
    channel = x.shape[-1]
    avg_pool = GlobalAveragePooling2D()(x)
    avg_pool = Dense(channel // ratio, activation='relu')(avg_pool)
    avg_pool = Dense(channel, activation='sigmoid')(avg_pool)

    return Multiply()([x, avg_pool])

# import tensorflow.keras.backend as K
# from tensorflow.keras.layers import Conv2D

from tensorflow.keras.layers import Conv2D, Lambda

# Modified SpatialAttention function using Keras layers only
def SpatialAttention(x):
    # Calculate average and max pooling along the channel axis, using Keras
    ↪layers
    avg_pool = Lambda(lambda y: K.mean(y, axis=-1, keepdims=True))(x)
    max_pool = Lambda(lambda y: K.max(y, axis=-1, keepdims=True))(x)

    # Concatenate the pooled features along the channel axis
    concat = concatenate([avg_pool, max_pool], axis=-1)

    # Convolution layer to create the attention map
    attention = Conv2D(1, (7, 7), padding="same", activation="sigmoid")(concat)

    # Multiply attention map with the input feature map
    return tf.keras.layers.multiply([x, attention])

# SDI (Semantics and Detail Infusion) Module

```

```

from tensorflow.keras.layers import UpSampling2D, MaxPooling2D, Conv2D

# Updated SDI (Semantics and Detail Infusion) Module with matching dimensions
def SDI_module(x):
    x_upsample = UpSampling2D()(x) # Upsampling the feature maps
    x_downsample = MaxPooling2D()(x) # Downsampling the feature maps
    identity_map = x # Identity map (preserve original detail)

    # Resize upsampled and downsampled features to match identity_map dimensions
    x_upsample = Conv2D(x.shape[-1], (1, 1),
        ↪padding='same')(UpSampling2D(size=(2, 2))(x_downsample)) # Match dimensions
    x_downsample = Conv2D(x.shape[-1], (1, 1),
        ↪padding='same')(UpSampling2D(size=(2, 2))(x_downsample)) # Match dimensions

    # Combine all three maps with attention mechanisms
    fused = concatenate([x_upsample, x_downsample, identity_map], axis=-1)
    fused = ChannelAttention(fused)
    fused = SpatialAttention(fused)

    return fused

# SmoothConv: A smoothing convolutional layer
def SmoothConv(x):
    return Conv2D(filters=x.shape[-1], kernel_size=(3, 3), padding="same")(x)

# U-Net v2 Architecture
def unet_v2(input_size=(256, 256, 3)):
    inputs = Input(input_size)

    # First DownConvolution / Encoder Leg
    conv1 = Conv2D(64, (3, 3), padding="same")(inputs)
    bn1 = Activation("relu")(conv1)
    conv1 = Conv2D(64, (3, 3), padding="same")(bn1)
    bn1 = BatchNormalization(axis=3)(conv1)
    bn1 = Activation("relu")(bn1)
    pool1 = MaxPooling2D(pool_size=(2, 2))(bn1)

    conv2 = Conv2D(128, (3, 3), padding="same")(pool1)
    bn2 = Activation("relu")(conv2)
    conv2 = Conv2D(128, (3, 3), padding="same")(bn2)
    bn2 = BatchNormalization(axis=3)(conv2)
    bn2 = Activation("relu")(bn2)
    pool2 = MaxPooling2D(pool_size=(2, 2))(bn2)

    conv3 = Conv2D(256, (3, 3), padding="same")(pool2)
    bn3 = Activation("relu")(conv3)

```

```

conv3 = Conv2D(256, (3, 3), padding="same")(bn3)
bn3 = BatchNormalization(axis=3)(conv3)
bn3 = Activation("relu")(bn3)
pool3 = MaxPooling2D(pool_size=(2, 2))(bn3)

conv4 = Conv2D(512, (3, 3), padding="same")(pool3)
bn4 = Activation("relu")(conv4)
conv4 = Conv2D(512, (3, 3), padding="same")(bn4)
bn4 = BatchNormalization(axis=3)(conv4)
bn4 = Activation("relu")(bn4)
pool4 = MaxPooling2D(pool_size=(2, 2))(bn4)

# Bottleneck
conv5 = Conv2D(1024, (3, 3), padding="same")(pool4)
bn5 = Activation("relu")(conv5)
conv5 = Conv2D(1024, (3, 3), padding="same")(bn5)
bn5 = BatchNormalization(axis=3)(conv5)
bn5 = Activation("relu")(bn5)

# Decoder Leg / UpConvolution with SDI and Attention
up6 = concatenate([Conv2DTranspose(512, (2, 2), strides=(2, 2),
padding="same")(bn5), SDI_module(conv4)], axis=3)
conv6 = SmoothConv(up6)
bn6 = Activation("relu")(conv6)
conv6 = Conv2D(512, (3, 3), padding="same")(bn6)
bn6 = BatchNormalization(axis=3)(conv6)
bn6 = Activation("relu")(bn6)

up7 = concatenate([Conv2DTranspose(256, (2, 2), strides=(2, 2),
padding="same")(bn6), SDI_module(conv3)], axis=3)
conv7 = SmoothConv(up7)
bn7 = Activation("relu")(conv7)
conv7 = Conv2D(256, (3, 3), padding="same")(bn7)
bn7 = BatchNormalization(axis=3)(conv7)
bn7 = Activation("relu")(bn7)

up8 = concatenate([Conv2DTranspose(128, (2, 2), strides=(2, 2),
padding="same")(bn7), SDI_module(conv2)], axis=3)
conv8 = SmoothConv(up8)
bn8 = Activation("relu")(conv8)
conv8 = Conv2D(128, (3, 3), padding="same")(bn8)
bn8 = BatchNormalization(axis=3)(conv8)
bn8 = Activation("relu")(bn8)

up9 = concatenate([Conv2DTranspose(64, (2, 2), strides=(2, 2),
padding="same")(bn8), SDI_module(conv1)], axis=3)
conv9 = SmoothConv(up9)

```

```

bn9 = Activation("relu")(conv9)
conv9 = Conv2D(64, (3, 3), padding="same")(bn9)
bn9 = BatchNormalization(axis=3)(conv9)
bn9 = Activation("relu")(bn9)

# Final Layer
conv10 = Conv2D(1, (1, 1), activation="sigmoid")(bn9)

return Model(inputs=[inputs], outputs=[conv10])

```

DICE COEFFICIENT

```

[12]: # function to create dice coefficient
def dice_coef(y_true, y_pred, smooth=100):
    y_true_flatten = K.flatten(y_true)
    y_pred_flatten = K.flatten(y_pred)

    intersection = K.sum(y_true_flatten * y_pred_flatten)
    union = K.sum(y_true_flatten) + K.sum(y_pred_flatten)
    return (2 * intersection + smooth) / (union + smooth)

# function to create dice loss
def dice_loss(y_true, y_pred, smooth=100):
    return -dice_coef(y_true, y_pred, smooth)

```

PLOTTING

```

[13]: def show_images(images, masks):
    plt.figure(figsize=(12, 12))
    for i in range(25):
        plt.subplot(5, 5, i+1)
        img_path = images[i]
        mask_path = masks[i]
        # read image and convert it to RGB scale
        image = cv2.imread(img_path)
        image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
        # read mask
        mask = cv2.imread(mask_path)
        # sho image and mask
        plt.imshow(image)
        plt.imshow(mask, alpha=0.4)

        plt.axis('off')

    plt.tight_layout()
    plt.show()

```

DATASET


```
[14]: data_dir = '/kaggle/input/lgg-mri-segmentation/kaggle_3m'

df = create_df(data_dir)
train_df, valid_df, test_df = split_df(df)

tr_aug_dict = dict(rotation_range=0.2,
                    width_shift_range=0.05,
                    height_shift_range=0.05,
                    shear_range=0.05,
                    zoom_range=0.05,
                    horizontal_flip=True,
                    fill_mode='nearest')

train_gen = create_gens(train_df, aug_dict=tr_aug_dict)
valid_gen = create_gens(valid_df, aug_dict={})
test_gen = create_gens(test_df, aug_dict={})

show_images(list(train_df['images_paths']), list(train_df['masks_paths']))
```



MODEL SUMMARY

```
[15]: model = unet()
model.compile(Adamax(learning_rate= 0.001), loss= dice_loss, metrics= [
    ↪dice_coef])

model.summary()
```

Model: "functional_1"

Layer (type)	Output Shape	Param #	Connected to
--------------	--------------	---------	--------------

input_layer (InputLayer)	(None, 256, 256, 3)	0	-
conv2d (Conv2D)	(None, 256, 256, 64)	1,792	input_layer[0][0]
activation (Activation)	(None, 256, 256, 64)	0	conv2d[0][0]
conv2d_1 (Conv2D)	(None, 256, 256, 64)	36,928	activation[0][0]
batch_normalization (BatchNormalizatio...	(None, 256, 256, 64)	256	conv2d_1[0][0]
activation_1 (Activation)	(None, 256, 256, 64)	0	batch_normalizat...
max_pooling2d (MaxPooling2D)	(None, 128, 128, 64)	0	activation_1[0][...
conv2d_2 (Conv2D)	(None, 128, 128, 128)	73,856	max_pooling2d[0]...
activation_2 (Activation)	(None, 128, 128, 128)	0	conv2d_2[0][0]
conv2d_3 (Conv2D)	(None, 128, 128, 128)	147,584	activation_2[0][...
batch_normalizatio... (BatchNormalizatio...	(None, 128, 128, 128)	512	conv2d_3[0][0]
activation_3 (Activation)	(None, 128, 128, 128)	0	batch_normalizat...
max_pooling2d_1 (MaxPooling2D)	(None, 64, 64, 128)	0	activation_3[0][...
conv2d_4 (Conv2D)	(None, 64, 64, 256)	295,168	max_pooling2d_1[...
activation_4 (Activation)	(None, 64, 64, 256)	0	conv2d_4[0][0]
conv2d_5 (Conv2D)	(None, 64, 64, 256)	590,080	activation_4[0][...

batch_normalizatio... (BatchNormalizatio...	(None, 64, 64, 256)	1,024	conv2d_5[0][0]
activation_5 (Activation)	(None, 64, 64, 256)	0	batch_normalizat...
max_pooling2d_2 (MaxPooling2D)	(None, 32, 32, 256)	0	activation_5[0][...
conv2d_6 (Conv2D)	(None, 32, 32, 512)	1,180,160	max_pooling2d_2[...
activation_6 (Activation)	(None, 32, 32, 512)	0	conv2d_6[0][0]
conv2d_7 (Conv2D)	(None, 32, 32, 512)	2,359,808	activation_6[0][...
batch_normalizatio... (BatchNormalizatio...	(None, 32, 32, 512)	2,048	conv2d_7[0][0]
activation_7 (Activation)	(None, 32, 32, 512)	0	batch_normalizat...
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 512)	0	activation_7[0][...
conv2d_8 (Conv2D)	(None, 16, 16, 1024)	4,719,616	max_pooling2d_3[...
activation_8 (Activation)	(None, 16, 16, 1024)	0	conv2d_8[0][0]
conv2d_9 (Conv2D)	(None, 16, 16, 1024)	9,438,208	activation_8[0][...
batch_normalizatio... (BatchNormalizatio...	(None, 16, 16, 1024)	4,096	conv2d_9[0][0]
activation_9 (Activation)	(None, 16, 16, 1024)	0	batch_normalizat...
conv2d_transpose (Conv2DTranspose)	(None, 32, 32, 512)	2,097,664	activation_9[0][...
concatenate (Concatenate)	(None, 32, 32, 1024)	0	conv2d_transpose... conv2d_7[0][0]

conv2d_10 (Conv2D)	(None, 32, 32, 512)	4,719,104	concatenate[0][0]
activation_10 (Activation)	(None, 32, 32, 512)	0	conv2d_10[0][0]
conv2d_11 (Conv2D)	(None, 32, 32, 512)	2,359,808	activation_10[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 32, 32, 512)	2,048	conv2d_11[0][0]
activation_11 (Activation)	(None, 32, 32, 512)	0	batch_normalizat...
conv2d_transpose_1 (Conv2DTranspose)	(None, 64, 64, 256)	524,544	activation_11[0]...
concatenate_1 (Concatenate)	(None, 64, 64, 512)	0	conv2d_transpose... conv2d_5[0][0]
conv2d_12 (Conv2D)	(None, 64, 64, 256)	1,179,904	concatenate_1[0]...
activation_12 (Activation)	(None, 64, 64, 256)	0	conv2d_12[0][0]
conv2d_13 (Conv2D)	(None, 64, 64, 256)	590,080	activation_12[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 64, 64, 256)	1,024	conv2d_13[0][0]
activation_13 (Activation)	(None, 64, 64, 256)	0	batch_normalizat...
conv2d_transpose_2 (Conv2DTranspose)	(None, 128, 128, 128)	131,200	activation_13[0]...
concatenate_2 (Concatenate)	(None, 128, 128, 256)	0	conv2d_transpose... conv2d_3[0][0]
conv2d_14 (Conv2D)	(None, 128, 128, 128)	295,040	concatenate_2[0]...
activation_14 (Activation)	(None, 128, 128, 128)	0	conv2d_14[0][0]

conv2d_15 (Conv2D)	(None, 128, 128, 128)	147,584	activation_14[0]...
batch_normalizatio...	(None, 128, 128, 128)	512	conv2d_15[0][0]
activation_15	(None, 128, 128, 128)	0	batch_normalizat...
conv2d_transpose_3	(None, 256, 256, 64)	32,832	activation_15[0]...
concatenate_3	(None, 256, 256, 128)	0	conv2d_transpose...
conv2d_16 (Conv2D)	(None, 256, 256, 64)	73,792	concatenate_3[0]...
activation_16	(None, 256, 256, 64)	0	conv2d_16[0][0]
conv2d_17 (Conv2D)	(None, 256, 256, 64)	36,928	activation_16[0]...
batch_normalizatio...	(None, 256, 256, 64)	256	conv2d_17[0][0]
activation_17	(None, 256, 256, 64)	0	batch_normalizat...
conv2d_18 (Conv2D)	(None, 256, 256, 1)	65	activation_17[0]...

Total params: 31,043,521 (118.42 MB)

Trainable params: 31,037,633 (118.40 MB)

Non-trainable params: 5,888 (23.00 KB)

TRAINING

```
[16]: import math
from keras.callbacks import ModelCheckpoint

epochs = 25
```

```

batch_size = 40
callbacks = [ModelCheckpoint('/kaggle/working/unet.keras', verbose=0,
↪save_best_only=True)]

history = model.fit(train_gen,
                    steps_per_epoch=math.ceil(len(train_df) / batch_size),
                    epochs=epochs,
                    verbose=1,
                    callbacks=callbacks,
                    validation_data=valid_gen,
                    validation_steps=math.ceil(len(valid_df) / batch_size))

```

Found 3143 validated image filenames.

Found 3143 validated image filenames.

Epoch 1/25

2024-03-27 03:26:47.480555: E

external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm eng4{k11=2} for conv (f32[40,128,128,128]{3,2,1,0}, u8[0]{0}) custom-call(f32[40,128,128,128]{3,2,1,0}, f32[128,128,3,3]{3,2,1,0}), window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01, custom_call_target="__cudnn\$convBackwardInput", backend_config={"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is taking a while...

2024-03-27 03:26:47.574930: E

external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took 1.094569919s

Trying algorithm eng4{k11=2} for conv (f32[40,128,128,128]{3,2,1,0}, u8[0]{0}) custom-call(f32[40,128,128,128]{3,2,1,0}, f32[128,128,3,3]{3,2,1,0}), window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01, custom_call_target="__cudnn\$convBackwardInput", backend_config={"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is taking a while...

2024-03-27 03:28:06.727656: E

external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm eng0{} for conv (f32[64,64,3,3]{3,2,1,0}, u8[0]{0}) custom-call(f32[40,64,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}), window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01, custom_call_target="__cudnn\$convBackwardFilter", backend_config={"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is taking a while...

2024-03-27 03:28:12.031047: E

external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took 6.303566872s

Trying algorithm eng0{} for conv (f32[64,64,3,3]{3,2,1,0}, u8[0]{0}) custom-call(f32[40,64,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}), window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01, custom_call_target="__cudnn\$convBackwardFilter", backend_config={"conv_result_sc

```

ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:28:17.098433: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng0{} for conv (f32[64,64,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,64,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}), window={size=3x3
pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:28:22.444199: E
external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
6.345852364s
Trying algorithm eng0{} for conv (f32[64,64,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,64,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}), window={size=3x3
pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:29:50.437931: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng0{} for conv (f32[128,64,2,2]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,64,256,256]{3,2,1,0}, f32[40,128,128,128]{3,2,1,0}),
window={size=2x2 stride=2x2}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:29:51.082708: E
external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
1.644891291s
Trying algorithm eng0{} for conv (f32[128,64,2,2]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,64,256,256]{3,2,1,0}, f32[40,128,128,128]{3,2,1,0}),
window={size=2x2 stride=2x2}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:29:52.909779: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng0{} for conv (f32[128,64,2,2]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,64,256,256]{3,2,1,0}, f32[40,128,128,128]{3,2,1,0}),
window={size=2x2 stride=2x2}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:29:54.028149: E
external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
2.118466171s
Trying algorithm eng0{} for conv (f32[128,64,2,2]{3,2,1,0}, u8[0]{0}) custom-

```

```

call(f32[40,64,256,256]{3,2,1,0}, f32[40,128,128,128]{3,2,1,0}),
window={size=2x2 stride=2x2}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:01.481301: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng20{k2=1,k3=0} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:01.486867: E
external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
1.005676733s
Trying algorithm eng20{k2=1,k3=0} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0})
custom-call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:02.487052: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng1{k2=2,k3=0} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:02.550301: E
external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
1.063361436s
Trying algorithm eng1{k2=2,k3=0} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0})
custom-call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:03.550495: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng0{} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_scale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:15.297811: E

```

```

external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
12.747427894s
Trying algorithm eng0{} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:22.232579: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng20{k2=1,k3=0} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:22.246161: E
external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
1.0137615s
Trying algorithm eng20{k2=1,k3=0} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0})
custom-call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:23.246413: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng1{k2=2,k3=0} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:23.316298: E
external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
1.070058532s
Trying algorithm eng1{k2=2,k3=0} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0})
custom-call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:24.316554: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng0{} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc

```



```

ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:30:36.557383: E
external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
13.241000245s
Trying algorithm eng0{} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[40,128,256,256]{3,2,1,0}, f32[40,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
WARNING: All log messages before absl::InitializeLog() is called are written to
STDERR
W0000 00:00:1711510238.150314      89 hlo_rematerialization.cc:2946] Can't
reduce memory use below 10.84GiB (11641402856 bytes) by rematerialization; only
reduced to 11.76GiB (12623817276 bytes), down from 14.33GiB (15388767268 bytes)
originally
I0000 00:00:1711510243.181085      89 device_compiler.h:186] Compiled cluster
using XLA! This line is logged at most once for the lifetime of the process.

78/79          1s 2s/step -
dice_coef: 0.1071 - loss: -0.1071

2024-03-27 03:35:35.161772: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng0{} for conv (f32[64,64,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[23,64,256,256]{3,2,1,0}, f32[23,64,256,256]{3,2,1,0}), window={size=3x3
pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:35:37.510867: E
external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
3.349190802s
Trying algorithm eng0{} for conv (f32[64,64,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[23,64,256,256]{3,2,1,0}, f32[23,64,256,256]{3,2,1,0}), window={size=3x3
pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:35:40.969204: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng0{} for conv (f32[64,64,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[23,64,256,256]{3,2,1,0}, f32[23,64,256,256]{3,2,1,0}), window={size=3x3
pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:35:43.314237: E

```

```

external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
3.345188288s
Trying algorithm eng0{} for conv (f32[64,64,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[23,64,256,256]{3,2,1,0}, f32[23,64,256,256]{3,2,1,0}), window={size=3x3
pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:36:43.251431: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng0{} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[23,128,256,256]{3,2,1,0}, f32[23,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:36:48.895652: E
external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
6.644334534s
Trying algorithm eng0{} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[23,128,256,256]{3,2,1,0}, f32[23,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:36:54.420730: E
external/local_xla/xla/service/slow_operation_alarm.cc:65] Trying algorithm
eng0{} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[23,128,256,256]{3,2,1,0}, f32[23,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...
2024-03-27 03:37:00.058052: E
external/local_xla/xla/service/slow_operation_alarm.cc:133] The operation took
6.637428608s
Trying algorithm eng0{} for conv (f32[64,128,3,3]{3,2,1,0}, u8[0]{0}) custom-
call(f32[23,128,256,256]{3,2,1,0}, f32[23,64,256,256]{3,2,1,0}),
window={size=3x3 pad=1_1x1_1}, dim_labels=bf01_oi01->bf01,
custom_call_target="__cudnn$convBackwardFilter", backend_config={"conv_result_sc
ale":1,"activation_mode":"kNone","side_input_scale":0,"leakyrelu_alpha":0} is
taking a while...

79/79                0s 5s/step -
dice_coef: 0.1079 - loss: -0.1079 Found 393 validated image filenames.
Found 393 validated image filenames.
79/79                886s 6s/step -
dice_coef: 0.1087 - loss: -0.1087 - val_dice_coef: 0.0207 - val_loss: -0.0207

```

Epoch 2/25
79/79 160s 2s/step -
dice_coef: 0.3575 - loss: -0.3575 - val_dice_coef: 0.0108 - val_loss: -0.0108
Epoch 3/25
79/79 162s 2s/step -
dice_coef: 0.5777 - loss: -0.5777 - val_dice_coef: 0.2745 - val_loss: -0.2745
Epoch 4/25
79/79 161s 2s/step -
dice_coef: 0.6652 - loss: -0.6652 - val_dice_coef: 0.4538 - val_loss: -0.4538
Epoch 5/25
79/79 161s 2s/step -
dice_coef: 0.7094 - loss: -0.7094 - val_dice_coef: 0.5693 - val_loss: -0.5693
Epoch 6/25
79/79 161s 2s/step -
dice_coef: 0.7283 - loss: -0.7283 - val_dice_coef: 0.6798 - val_loss: -0.6798
Epoch 7/25
79/79 161s 2s/step -
dice_coef: 0.7370 - loss: -0.7370 - val_dice_coef: 0.7237 - val_loss: -0.7237
Epoch 8/25
79/79 161s 2s/step -
dice_coef: 0.7572 - loss: -0.7572 - val_dice_coef: 0.7285 - val_loss: -0.7285
Epoch 9/25
79/79 160s 2s/step -
dice_coef: 0.7747 - loss: -0.7747 - val_dice_coef: 0.7063 - val_loss: -0.7063
Epoch 10/25
79/79 160s 2s/step -
dice_coef: 0.7748 - loss: -0.7748 - val_dice_coef: 0.6966 - val_loss: -0.6966
Epoch 11/25
79/79 161s 2s/step -
dice_coef: 0.8012 - loss: -0.8012 - val_dice_coef: 0.7816 - val_loss: -0.7816
Epoch 12/25
79/79 159s 2s/step -
dice_coef: 0.7868 - loss: -0.7868 - val_dice_coef: 0.7721 - val_loss: -0.7721
Epoch 13/25
79/79 160s 2s/step -
dice_coef: 0.7850 - loss: -0.7850 - val_dice_coef: 0.7731 - val_loss: -0.7731
Epoch 14/25
79/79 161s 2s/step -
dice_coef: 0.8147 - loss: -0.8147 - val_dice_coef: 0.7894 - val_loss: -0.7894
Epoch 15/25
79/79 160s 2s/step -
dice_coef: 0.8198 - loss: -0.8198 - val_dice_coef: 0.6664 - val_loss: -0.6664
Epoch 16/25
79/79 159s 2s/step -
dice_coef: 0.8315 - loss: -0.8315 - val_dice_coef: 0.7883 - val_loss: -0.7883
Epoch 17/25
79/79 160s 2s/step -
dice_coef: 0.8233 - loss: -0.8233 - val_dice_coef: 0.7739 - val_loss: -0.7739

```

Epoch 18/25
79/79          159s 2s/step -
dice_coef: 0.8205 - loss: -0.8205 - val_dice_coef: 0.7775 - val_loss: -0.7775
Epoch 19/25
79/79          159s 2s/step -
dice_coef: 0.8255 - loss: -0.8255 - val_dice_coef: 0.7525 - val_loss: -0.7525
Epoch 20/25
79/79          160s 2s/step -
dice_coef: 0.8218 - loss: -0.8218 - val_dice_coef: 0.7826 - val_loss: -0.7826
Epoch 21/25
79/79          161s 2s/step -
dice_coef: 0.8428 - loss: -0.8428 - val_dice_coef: 0.8089 - val_loss: -0.8089
Epoch 22/25
79/79          161s 2s/step -
dice_coef: 0.8481 - loss: -0.8481 - val_dice_coef: 0.8360 - val_loss: -0.8360
Epoch 23/25
79/79          159s 2s/step -
dice_coef: 0.8420 - loss: -0.8420 - val_dice_coef: 0.7910 - val_loss: -0.7910
Epoch 24/25
79/79          159s 2s/step -
dice_coef: 0.8432 - loss: -0.8432 - val_dice_coef: 0.8328 - val_loss: -0.8328
Epoch 25/25
79/79          161s 2s/step -
dice_coef: 0.8712 - loss: -0.8712 - val_dice_coef: 0.8585 - val_loss: -0.8585

```

```
[17]: model.save("/kaggle/working/model.h5")
```

LOAD MODAL

```
[ ]: from keras.models import load_model

# Assuming unet.keras is located in the current working directory
model = load_model('/kaggle/working/unet.keras')

# Now you can use the loaded model for inference or further processing

```

DICE SCORE COMPUTATION

```
[ ]: ts_length = len(test_df)
test_batch_size = max(sorted([ts_length // n for n in range(1, ts_length + 1)]
    ↪ if ts_length%n == 0 and ts_length/n <= 80]))
test_steps = ts_length // test_batch_size

train_score = model.evaluate(train_gen, steps= test_steps, verbose= 1)
valid_score = model.evaluate(valid_gen, steps= test_steps, verbose= 1)
test_score = model.evaluate(test_gen, steps= test_steps, verbose= 1)

```

DICE VALUE

```
[1]: print("Train Loss: ", train_score[0])
      print("Train Dice: ", train_score[1])
      print('-' * 20)

      print("Valid Loss: ", valid_score[0])
      print("Valid Dice: ", valid_score[1])
      print('-' * 20)

      print("Test Loss: ", test_score[0])
      print("Test Dice: ", test_score[1])
```

```
Training Loss: 0.1224716271181162
Training Dice: 0.8721450567245488
-----
Validation Loss: 0.1395732432585654
Validation Dice: 0.8610987663269046
-----
Test Loss: 0.130535134252623
Test Dice: 0.870535135269167
```

PLOTTING

```
[27]: import numpy as np
      import matplotlib.pyplot as plt

      def plot_training(hist):
          """
          This function takes a training model and plots the history of accuracy and
          losses with the best epoch in both of them.
          """

          # Define needed variables
          tr_dice = hist.history['dice_coef']
          tr_loss = hist.history['loss']

          val_dice = hist.history['val_dice_coef']
          val_loss = hist.history['val_loss']

          index_dice = np.argmax(val_dice) # Corrected line
          dice_highest = val_dice[index_dice]
          index_loss = np.argmin(val_loss)
          val_lowest = val_loss[index_loss]

          Epochs = [i+1 for i in range(len(val_loss))]

          dice_label = f'best epoch= {str(index_dice + 1)}'
          loss_label = f'best epoch= {str(index_loss + 1)}'
```

```

# Plot training history
plt.figure(figsize=(20, 20))
plt.style.use('fivethirtyeight')


# Training Dice
plt.subplot(2, 2, 3)
plt.plot(Epochs, tr_dice, 'r', label='Training Dice')
plt.plot(Epochs, val_dice, 'g', label='Validation Dice')
plt.scatter(index_dice + 1, dice_highest, s=150, c='blue',
↪label=dice_label)
plt.title('Training and Validation Dice Coefficient')
plt.xlabel('Epochs')
plt.ylabel('Dice')
plt.legend()

# Training Loss
plt.subplot(2, 2, 4)
plt.plot(Epochs, tr_loss, 'r', label='Training loss')
plt.plot(Epochs, val_loss, 'g', label='Validation loss')
plt.scatter(index_loss + 1, val_lowest, s=150, c='blue', label=loss_label)
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()

```

[28]: plot_training(history)



OUTPUT

```
[20]: for _ in range(20):
        index = np.random.randint(1, len(test_df.index))
        img = cv2.imread(test_df['images_paths'].iloc[index])
        img = cv2.resize(img, (256, 256))
        img = img/255
        img = img[np.newaxis, :, :, :]

        predicted_img = model.predict(img)

        plt.figure(figsize=(12, 12))


        plt.subplot(1, 3, 1)
        plt.imshow(np.squeeze(img))
        plt.axis('off')
        plt.title('Original Image')

        plt.subplot(1, 3, 2)
        plt.imshow(np.squeeze(cv2.imread(test_df['masks_paths'].iloc[index])))
        plt.axis('off')
        plt.title('Original Mask')

        plt.subplot(1, 3, 3)
        plt.imshow(np.squeeze(predicted_img) > 0.5 )
        plt.title('Prediction')
        plt.axis('off')

        plt.show()
```


1/1 9s 9s/step



1/1

0s 19ms/step


Original Image



Original Mask




Prediction




1/1

0s 19ms/step

Original Image



Original Mask




Prediction



1/1

0s 19ms/step

Original Image



Original Mask




Prediction




1/1

0s 19ms/step

Original Image



Original Mask




Prediction




1/1

0s 22ms/step


Original Image



Original Mask




Prediction



1/1

0s 22ms/step

Original Image



Original Mask




Prediction




1/1


0s 20ms/step




1/1 0s 19ms/step




1/1 0s 18ms/step



1/1 0s 20ms/step



1/1 0s 20ms/step




1/1 0s 19ms/step




1/1 0s 18ms/step

Original Image



Original Mask




Prediction




1/1

0s 19ms/step

Original Image



Original Mask




Prediction



1/1

0s 18ms/step


Original Image



Original Mask




Prediction




1/1

0s 19ms/step


Original Image



Original Mask




Prediction




1/1

0s 18ms/step

Original Image



Original Mask




Prediction



1/1

0s 18ms/step

Original Image



Original Mask




Prediction




1/1

0s 19ms/step

Original Image



Original Mask



Prediction



[]: